
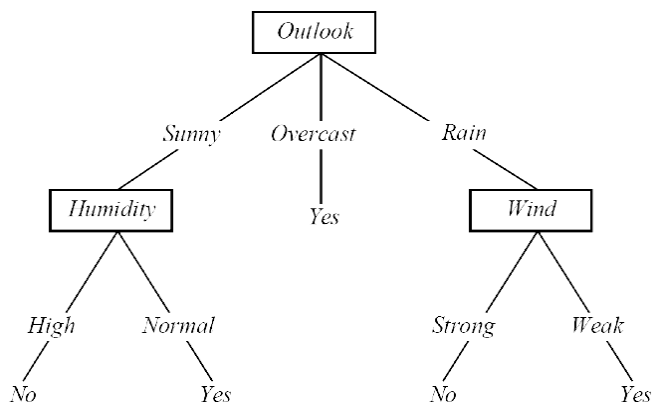


Learning

We have seen machine learning with different representations:


- (1) Decision trees -- symbolic representation of various decision rules -- “disjunction of conjunctions”
- (2) Perceptron -- learning of weights that represent a linear decision surface classifying a set of objects into two groups



Different representations give rise to different hypothesis or model spaces. Machine learning algorithms search these model spaces for the best fitting model.


Perceptron Learning Revisited

```
Initialize  $\bar{w}$  and  $b$  to random values.  
repeat  
  for each  $(\bar{x}_i, y_i) \in D$  do  
    if  $\hat{f}(\bar{x}_i) \neq y_i$  then  
      Update  $\bar{w}$  and  $b$  incrementally.  
    end if  
  end for  
until  $D$  is perfectly classified.  
return  $\bar{w}$  and  $b$ 
```




Constructs a line (hyperplane) as a classifier.

What About Non-Linearity?

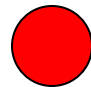



Can we learn this decision surface? ...Yes! Multi-Layer Perceptrons!

Multi-Layer Perceptrons



 \equiv Linear Unit

 \equiv



How do we train?

Perceptron was easy:

```
Initialize  $\bar{w}$  and  $b$  to random values.  
repeat  
  for each  $(\bar{x}_i, y_i) \in D$  do  
    if  $\hat{f}(\bar{x}_i) \neq y_i$  then  
      Update  $\bar{w}$  and  $b$  incrementally.  
    end if  
  end for  
until  $D$  is perfectly classified.  
return  $\bar{w}$  and  $b$ 
```

error




Every time we found an error of the predicted value $f(x_i)$ compared to the label in the training set y_i , we update w and b .

Artificial Neural Networks

Feed-forward with Backpropagation

We have to be a bit smarter in the case of ANNs: compute the error (feed forward) and then use the error to update all the weights by propagating the error back.



Backpropagation


$$\Delta = (t - y)^2$$

$$\delta_o = y(1 - y)\Delta$$

$$w_{ho} \leftarrow w_{ho} + \alpha_o \delta_o$$

$$\delta_h = y(1 - y)w_{ho} \delta_o$$

$$w_{ih} \leftarrow w_{ih} + \alpha_h \delta_h$$



This only works because

$$\delta_o = y(1 - y)\Delta = \frac{\partial \Delta}{\partial w \bullet x} = \frac{\partial (t - y)^2}{\partial w \bullet x}$$

and the output y is differentiable because the transfer function is differentiable. Also note, everything is based on the *rate of change* of the error...we are searching in the direction where the rate of change will minimize the output error.

Backpropagation Algorithm

Note: this algorithm is for a NN with a single output node o and a single hidden layer. It can easily be generalized.

Initialize the weights in the network (often randomly)

Do

For each example e in the training set

// forward pass

y = compute neural net output

t = label for e from training data

Calculate error $\Delta = (t - y)^2$ at the output units

// backward pass

Compute error δ_o for weights from a hidden node h to the output node o using Δ

Compute error δ_h for weights from an input node i to hidden node h using δ_o

Update the weights in the network

Until all examples classified correctly or stopping criterion satisfied

Return the network

Neural Network Learning


- Define the network error as

$$\Delta_x = (t - y)^2$$

for some $x \in X$, where i is an index over the output units.



- Let $\Delta_x(w)$ be the error E_x as a function of the weights w .
- Use the gradient (slope) of the error surface to guide the search towards appropriate weights:

$$\Delta w_k = -\eta \frac{\partial \Delta_x}{\partial w_k}$$




Representational Power

- Every bounded continuous function can be approximated with arbitrarily small error by a network with one hidden layer.
- Any function can be approximated to arbitrary accuracy by a network with two hidden layers.



Hidden Layer Representations


Target Function:



Input	Output
10000000	→ 10000000
01000000	→ 01000000
00100000	→ 00100000
00010000	→ 00010000
00001000	→ 00001000
00000100	→ 00000100
00000010	→ 00000010
00000001	→ 00000001

Can this be learned?

Hidden Layer Representations



Input		Hidden Values		Output
10000000	→	.89 .04 .08	→	10000000
01000000	→	.01 .11 .88	→	01000000
00100000	→	.01 .97 .27	→	00100000
00010000	→	.99 .97 .71	→	00010000
00001000	→	.03 .05 .02	→	00001000
00000100	→	.22 .99 .99	→	00000100
00000010	→	.80 .01 .98	→	00000010
00000001	→	.60 .94 .01	→	00000001

1 0 0
0 0 1
0 1 0
1 1 1
0 0 0
0 1 1
1 0 1
1 1 0

Hidden layers allow a network to invent appropriate internal representations.

MLPClassifier

```
import pandas as pd
from sklearn.neural_network import MLPClassifier
from sklearn.model_selection import cross_val_score
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix

# get data
df = pd.read_csv("wdbc.csv")
df = df.drop(['ID'],axis=1)
X = df.drop(['Diagnosis'],axis=1)
y = df['Diagnosis']

# neural network
model = MLPClassifier(hidden_layer_sizes=(15,))

# do the 5-fold cross validation
scores = cross_val_score(model, X, y, cv=5)
print("Fold Accuracies: {}".format(scores))
print("Accuracy: {}".format(scores.mean()))
```

Fold Accuracies: [0.87591241
0.95620438 0.95620438 0.98540146
0.98518519]

Accuracy: 0.9517815625844823

MLP Grid Search

We can also perform a grid search

BEWARE: a grid search over all possible parameters of an MLP is almost impossible - combinatoric explosion, too many different combinations possible.

Here we only perform a grid over the *number of nodes in a single hidden layer*.

MLP Grid Search

```
# neural network
```

```
model = MLPClassifier(max_iter=2000)
```

```
# grid search
```

```
param_grid = {'hidden_layer_sizes': [(5,), (6,), (7,), (8,), (9,), (10,),  
                                     (11,), (12,), (13,), (14,), (15,), (16,),  
                                     (17,), (18,), (19,), (20,)]}
```

```
grid = GridSearchCV(model, param_grid, cv=5)
```

```
grid.fit(X, actual_y)
```

```
print("Grid Search: best parameters: {}".format(grid.best_params_))
```

```
# evaluate the best model
```

```
best_model = grid.best_estimator_
```

```
predict_y = best_model.predict(X)
```

```
print("Accuracy: {}".format(accuracy_score(actual_y, predict_y)))
```

```
# build the confusion matrix
```

```
labels = ['benign', 'malignant']
```

```
cm = confusion_matrix(actual_y, predict_y, labels=labels)
```

```
cm_df = pd.DataFrame(cm, index=labels, columns=labels)
```

```
print("Confusion Matrix:\n{}".format(cm_df))
```

```
# bootstrapped confidence interval
```

```
print("Confidence interval best MLP: {}".format(bootstrap(best_model, df, 'class')))
```

MLP Grid Search

```
Grid Search: best parameters: {'hidden_layer_sizes': (9,)}
```

```
Accuracy: 0.9707174231332357
```

```
Confusion Matrix:
```

	benign	malignant
benign	435	9
malignant	11	228

```
Confidence interval best MLP: (0.93412408759124088, 0.9928832116788322)
```


Team Exercise

Use the Crohn's Disease dataset: CrohnD

<https://vincentarelbundock.github.io/Rdatasets/datasets.html>

You will need to preprocess this before you can use it:

c1 -> 0, c2 -> 1, F -> 0, M -> 1

Build a ANN/MLP with the best cross-validated performance you can find.

Compare it to either a tree or a KNN (or both).

Report if the difference between the models is statistically significant.

Teams:

Team 0: Shehjar Harout Geron

Team 1: Aakash Cory Christopher

Team 2: Maurice Kevin Ben

Team 3: Najib Aguilar Ronil

Team 4: Joe Shamal Gabe

Team 5: Kermalyn Peter David_M

Team 6: Matt Alexander Alber

Team 7: Evelyn Susallin David_P